

Assignment 1: The Causes of Economic Growth

Introduction to Applied Data Science

2022-2023

Bas Machielsen
a.h.machielsen@uu.nl

April 2023

Assignment 1: The Causes of Economic Growth

In this assignment, you will gather data from the *World Bank* website, and augment it with data from the *Clio Infra* website. Then, you will visualize these data using several graphs and tables, and test several hypotheses about the causes of economic growth. You will use this document to complete the code chunks which I have left unfinished to produce your own data analysis & visualization.

To start with, please replace my name and e-mail address with yours. Then, remove the lines:

```
output:
  pdf_document:
    includes:
      in_header: "preamble.tex"
```

from the document and replace them by:

```
output: pdf_document
```

Now, we're ready to start. For all of the code-related questions, please answer with code, and do not type (or copy) the answer from the console. Rather, let R *generate* your answer.

1. World Bank Data

The *World Bank* collects and processes large amounts of data and generates them on the basis of economic models. These data and models have gradually been made available to the public in a way that encourages reuse. In particular, the databases of the World Bank are available on <https://data.worldbank.org/>. It pays the effort to browse through the website, see if you can navigate your way through the website, and use the interface the World Bank provides you.

Normally, if you were looking for data from the World Bank, you would go to the website, find your dataset, download it to .xlsx or any other format, and then import it into an R `data.frame` using `read_xlsx()`, or something else. But, this reliance on manual downloads of spreadsheets of the data they are interested in can quickly become overwhelming, as the work is manual, time consuming, and not easily reproducible.

Fortunately, however, there also exist an R package which allows you to browse swiftly through World Bank data, and easily download it as an R `data.frame`. You can get this package by installing:

```
library(pacman)
p_load("wbstats", "tidyverse")
```

You can navigate the database by searching for terms:

```
wbstats::wb_search("gdp per capita")
```

```
## # A tibble: 24 x 3
##   indicator_id      indicator      indic-1
##   <chr>            <chr>            <chr>
## 1 5.51.01.10.gdp    Per capita GDP growth    GDP pe-
## 2 6.0.GDPpc_constant GDP per capita, PPP (constant 2011 international ~ GDP pe-
## 3 NV.AGR.PCAP.KD.ZG Real agricultural GDP per capita growth rate (%) The gr-
## 4 NY.GDP.PCAP.CD    GDP per capita (current US$)    GDP pe-
## 5 NY.GDP.PCAP.CN    GDP per capita (current LCU)    GDP pe-
## 6 NY.GDP.PCAP.KD    GDP per capita (constant 2010 US$)    GDP pe-
## 7 NY.GDP.PCAP.KD.ZG GDP per capita growth (annual %)    Annual~
## 8 NY.GDP.PCAP.KN    GDP per capita (constant LCU)    GDP pe-
## 9 NY.GDP.PCAP.PP.CD GDP per capita, PPP (current international $)    This i-
## 10 NY.GDP.PCAP.PP.KD GDP per capita, PPP (constant 2017 international ~ GDP pe-
## # ... with 14 more rows, and abbreviated variable name 1: indicator_desc
```

Afterwards, you can proceed to download data by executing `wb_data("indicator_id")`. You can then write this to a data.frame, and merge this data with other indicators to create a dataset. There exist many of these packages, and we will also use another today.

Apart from being easy to use, these packages also have another advantage: reproducibility. Collecting data by means of code allows other users to unambiguously reproduce your data collection process.

Firstly, we will look for GDP growth data.

Question x: pass a search query to `wb_search` for GDP growth data, and download the indicator for which the description matches “GDP (current US\$)”. The full description should read:

GDP at purchaser’s prices is the sum of gross value added by all resident producers in the economy plus any product taxes and minus any subsidies not included in the value of the products. It is calculated without making deductions for depreciation of fabricated assets or for depletion and degradation of natural resources. Data are in current U.S. dollars. Dollar figures for GDP are converted from domestic currencies using single year official exchange rates. For a few countries where the official exchange rate does not reflect the rate effectively applied to actual foreign exchange transactions, an alternative conversion factor is used.

```
wbstats::wb_search("GDP")
```

```
## # A tibble: 541 x 3
##   indicator_id      indicator      indic-1
##   <chr>            <chr>            <chr>
## 1 5.51.01.10.gdp    Per capita GDP growth    GDP pe-
## 2 6.0.GDP_current  GDP (current $)          GDP is-
## 3 6.0.GDP_growth   GDP growth (annual %)    Annual~
## 4 6.0.GDP_usd      GDP (constant 2005 $)    GDP is-
## 5 6.0.GDPpc_constant GDP per capita, PPP (constant 2011 internationa~ GDP pe-
## 6 BG.GSR.NFSV.GD.ZS Trade in services (% of GDP)    Trade ~
```

```
## 7 BI.WAG.TOTL.GD.ZS Wage bill as a percentage of GDP <NA>
## 8 BM.KLT.DINV.WD.GD.ZS Foreign direct investment, net outflows (% of G~ Foreign~
## 9 BN.CAB.XOKA.GD.ZS Current account balance (% of GDP) Curren~
## 10 BN.CUR.GDPM.ZS Current account balance excluding net official ~ Curren~
## # ... with 531 more rows, and abbreviated variable name 1: indicator_desc
```

```
gdp <- wbstats::wb_data('NY.GDP.MKTP.CD')
```

Question: Rename the variable NY.GDP.MKTP.CD to gdp. Remove the NA observations from the dataset. How many observations are there in the dataset in total?

```
gdp <- gdp %>%
  rename(gdp = `NY.GDP.MKTP.CD`) %>%
  filter(!is.na(gdp))
```

```
nrow(gdp)
```

```
## [1] 10336
```

Question x: How many observations per country are there? Show the first ten observations.

```
gdp %>%
  group_by(country) %>%
  summarize(count = n()) %>%
  head(10)
```

```
## # A tibble: 10 x 2
##   country      count
##   <chr>      <int>
## 1 Afghanistan    41
## 2 Albania        38
## 3 Algeria        62
## 4 American Samoa  19
## 5 Andorra        52
## 6 Angola         42
## 7 Antigua and Barbuda 45
## 8 Argentina      60
## 9 Armenia        32
## 10 Aruba         35
```

Question x: How many different years are there in the dataset? Put them in increasing order.

```
gdp %>%
  select(date) %>%
  pull() %>%
  unique() %>%
  sort()
```

```
## [1] 1960 1961 1962 1963 1964 1965 1966 1967 1968 1969 1970 1971 1972 1973 1974
## [16] 1975 1976 1977 1978 1979 1980 1981 1982 1983 1984 1985 1986 1987 1988 1989
## [31] 1990 1991 1992 1993 1994 1995 1996 1997 1998 1999 2000 2001 2002 2003 2004
## [46] 2005 2006 2007 2008 2009 2010 2011 2012 2013 2014 2015 2016 2017 2018 2019
## [61] 2020 2021
```

Question x: For each country, what is the first and last year? Again show the first ten observations.

```
gdp %>%
  group_by(country) %>%
  summarize(first_year = min(date), last_year = max(date))
```

```
## # A tibble: 214 x 3
##   country          first_year last_year
##   <chr>            <dbl>      <dbl>
## 1 Afghanistan      1960        2020
## 2 Albania           1984        2021
## 3 Algeria           1960        2021
## 4 American Samoa    2002        2020
## 5 Andorra           1970        2021
## 6 Angola            1980        2021
## 7 Antigua and Barbuda 1977        2021
## 8 Argentina         1962        2021
## 9 Armenia           1990        2021
## 10 Aruba            1986        2020
## # ... with 204 more rows
```

Question: Make a summary of the data, with the mean, median, sd, min and max values for gdp.

```
gdp %>%
  summarize(mean = mean(gdp),
            median = median(gdp),
            sd = sd(gdp),
            min = min(gdp),
            max = max(gdp))
```

```
## # A tibble: 1 x 5
##       mean      median      sd      min      max
##   <dbl>    <dbl>    <dbl>  <dbl>  <dbl>
## 1 184192121208. 7824737792. 991080605921. 8824448. 2.30e13
```

Question: What country, in which year, had the lowest GDP? And the highest? Hint: use a function similar to slice from the dplyr package.

```
slice_min(gdp, gdp, n = 1)
```

```
## # A tibble: 1 x 9
##   iso2c iso3c country  date      gdp unit  obs_status footnote last_updated
##   <chr> <chr> <chr>   <dbl>    <dbl> <chr> <chr>      <chr>      <date>
## 1 TV    TUV    Tuvalu  1990 8824448. <NA> <NA>      <NA>      2022-09-16
```

```
slice_max(gdp, gdp, n = 1)
```

```
## # A tibble: 1 x 9
##   iso2c iso3c country  date      gdp unit  obs_status footnote last_updated
##   <chr> <chr> <chr>   <dbl>    <dbl> <chr> <chr>      <chr>      <date>
## 1 US    USA    United States  2021 2.30e13 <NA> <NA>      <NA>      2022-09-16
```

Next, we'll have a look at population data, which we can also retrieve from the World Bank database.

```
population <- wb_data("SP.POP.TOTL")
```

Question: Rename the population variable to population and overwrite this to memory.

```
population <- population %>%  
  rename(population = `SP.POP.TOTL`)
```

Finally, we'll merge population with gdp on the basis of *country* and *year*.

Question: use `left_join` to merge gdp (left data.frame) with population (right data frame). Check whether everything has gone correctly. Save this dataframe to memory as `gdp_pop`. Select only *country*, *date*, *iso3c.x*, *gdp* and *population*. Then, use `mutate()` to create a new variable, `gdp_cap = gdp / population`. Then, again apply `na.omit()`. Write this data.frame to memory to `data`.

```
data <- left_join(gdp, population,  
  by = c("country" = "country",  
        "date" = "date")) %>%  
  select(country, date, iso3c.x, gdp, population) %>%  
  na.omit() %>%  
  mutate(gdp_cap = gdp / population)
```

Now, let's collect a pre-made version of GDP per capita from the World Bank website.

```
wb_search("gdp per capita")  
  
alt_gdp_pc <- wb_data('NY.GDP.PCAP.CD') %>%  
  filter(!is.na(`NY.GDP.PCAP.CD`))
```

Question: What is the correlation between these two variables? What does that mean?

```
cor(data$gdp_cap, alt_gdp_pc$NY.GDP.PCAP.CD)
```

2. World Tables: Capital Stock

Next, we'll proceed to find some potential determinants of GDP growth. One of the classical determinants of GDP per capita growth is the level of physical capital. Many models in macroeconomics explain economic well-being on account of the amount of capital in an economy. In particular, we'll look for a few measures from the Penn World Tables. This data has to be downloaded manually from [this website](#). You can either manually download an Excel file, in which case, make sure to put it in the right directory when reading it. Or, you can run the following chunk, which downloads the file to your working directory:

```
pacman::p_load("readxl")  
  
wd <- getwd()  
download.file("https://www.rug.nl/ggdc/docs/pwt100.xlsx", file.path(wd, "pwt100.xlsx"))  
  
pwt <- readxl::read_excel('pwt100.xlsx', sheet = 3)
```

We are looking for the *cn* variable, which indicates *Capital stock at current PPPs (in mil. 2017US\$)*.

Question: Select the variables *countrycode*, *year*, and *cn*, and rewrite the dataframe to memory.

```
pwt <- pwt %>%  
  select(countrycode, year, cn)
```

Now, we're looking to merge the two data.frames data and pwt, on the basis of common country names and years.

Question: Have a look at the two datasets below. On the basis of which two matched variables in both datasets do you have to perform the merge?

```
pwt
```

```
## # A tibble: 12,810 x 3
##   countrycode year   cn
##   <chr>      <dbl> <dbl>
## 1 ABW        1950    NA
## 2 ABW        1951    NA
## 3 ABW        1952    NA
## 4 ABW        1953    NA
## 5 ABW        1954    NA
## 6 ABW        1955    NA
## 7 ABW        1956    NA
## 8 ABW        1957    NA
## 9 ABW        1958    NA
## 10 ABW       1959    NA
## # ... with 12,800 more rows
```

```
data
```

```
## # A tibble: 10,333 x 6
##   country date iso3c.x      gdp population gdp_cap
##   <chr>   <dbl> <chr>      <dbl>      <dbl>    <dbl>
## 1 Aruba   1986 ABW      405586592.    62645    6474.
## 2 Aruba   1987 ABW      487709497.    61838    7887.
## 3 Aruba   1988 ABW      596648045.    61072    9770.
## 4 Aruba   1989 ABW      695530726.    61033   11396.
## 5 Aruba   1990 ABW      764804469.    62152   12305.
## 6 Aruba   1991 ABW      872067039.    64623   13495.
## 7 Aruba   1992 ABW      958659218.    68240   14048.
## 8 Aruba   1993 ABW     1083240223.    72495   14942.
## 9 Aruba   1994 ABW     1245810056.    76705   16242.
## 10 Aruba  1995 ABW     1320670391.    80324   16442.
## # ... with 10,323 more rows
```

We can do this in various ways: we can perform `left_join`, `right_join`, `inner_join`, or `outer_join`, but we can also use the `merge` function. Although their arguments differ somewhat, the results should absolutely be the same provided you specify the `by` arguments correctly. In that case, you match one particular country-year from the left data.frame to the identical particular country-year from the right data-frame, and put all variables together.

Question: Do this. Perform a merge and save the resulting data.frame to `merged_data`.

```
merged_data <- merge(data, pwt, by.x = c("iso3c.x", "date"), by.y=c("countrycode", "year"))
```

Next, we want to average GDP per capita and Capital stock for each country present in the dataset. This can be done easily using the `mutate` function from the `tidyverse` package. However, we want to investigate *current* GDP per capita, so we do not want to take too long an average.

Say we want to take an average over the years 2010-2020.

Question: Finish the following code to compute the average of GDP per capita and Capital stock. Make sure you deal with NA's explicitly. Save this again to merged_data

```
merged_data <- merged_data %>%
  group_by(country, iso3c.x) %>%
  filter(between(date, 2010, 2020)) %>%
  summarize(
    # fill in your answer here
    avg_gdpc = mean(gdp_cap, na.rm = TRUE),
    avg_cn = mean(cn, na.rm = TRUE)
  )
```

3. Historical Antecedents: Clio Infra Data

Another important element of capital might be not only physical capital, but human capital! Instead of using *contemporary* human capital to explain economic development in 2010-2020, we will use *historical* human capital. For this, we can again use a package, called Clio, which aggregates various historical datasets. You can install and load it by:

```
devtools::install_github("bas92/Clio")
library(Clio)
```

You can see what variables are available in this dataset by running:

```
Clio::clio_overview() %>% head(10)
```

```
##      variable_name from  to  obs
## 1  Cattle per Capita 1500 2010 7456
## 2  Cropland per Capita 1500 2010 6226
## 3  Goats per Capita 1500 2010 7037
## 4  Pasture per Capita 1500 2010 5963
## 5  Pigs per Capita 1500 2010 6841
## 6  Sheep per Capita 1500 2010 6835
## 7  Total Cattle 1500 2010 7457
## 8  Total Cropland 1500 2010 6191
## 9  Total Number of Goats 1500 2010 7037
## 10 Total Number of Pigs 1500 2010 6841
```

As a proxy to measure historical human capital, we'll use *Average Years of Education* in 1930. We can download this by running:

```
educ <- Clio::clio_get("Average Years of Education")
```

Question: Filter this dataset such that only observations from 1930 remain.

```
educ <- educ %>%
  filter(year == 1930)
```

Question: Merge this dataset to the merged_data set, so that the human capital proxy is added to the dataset.

When discussing economic growth, some people also talk about a *reversal of fortune* tendency: the countries that were relatively the most wealthy in or before the Middle Ages are among the poorest now, and vice versa.

	mean	median	min	max	N
avg_gdpc	15 798.78	6584.37	253.76	113 920.25	179
avg_cn	2 754 179.37	320 940.75	2556.14	75 221 956.00	176
av_educ	3.22	2.62	0.17	8.30	46
urban	0.09	0.08	0.00	0.32	46

We also want to investigate such an hypothesis. In order to do so, we use a proxy for wealth from 1500, the urbanization ratio. This can also be downloaded from the Clio Infra database.

Question: Now find and download Urbanization Ratio, filter the dataset such that only observations from 1500 remain, and save it to urb.

```
urb <- clio_get("Urbanization Ratio") %>%
  filter(year == 1500)
```

Now, let's merge educ and urb together, and then subsequently merge this to the merged_data data.frame.

```
educ_urb <- merge(educ, urb, by = "ccode")
```

Question: Also remove year.x, year.y and country.name.y from the dataset.

```
educ_urb <- educ_urb %>%
  select(-c(year.x, year.y, country.name.y))
```

Finally, we need to merge educ_urb to the merged_data data.frame.

Question: Merge these two dataframes using left_join, with merged_data being the left data.frame. Rename Average Years of Education and Urbanization Ratio to av_educ and urban respectively.

```
final <- merged_data %>%
  left_join(educ_urb,
            by = c("country" = "country.name.x")) %>%
  rename(av_educ = `Average Years of Education`, urban = `Urbanization Ratio`)
```

4. Summarizing and Analyzing the Data

Question: Create a descriptive statistics table using the variables we have obtained. In it, we want to display the mean, median, sd, min, max and number of observations. Hint: use the modelsummary package, and use the following syntax for the formula: x1 + x2 + x3 + x4 ~ (mean + median + min + max)*Arguments(na.rm = TRUE) + N.

```
library(modelsummary)

modelsummary::datasummary(data = final, avg_gdpc + avg_cn + av_educ + urban ~
  (mean + median + min + max)*Arguments(na.rm = TRUE) + N)
```


Next, we would like to make a histogram of the four aforementioned variables, to investigate their distribution. To do so, we need data in so-called long-form:

```
final_long <- final %>%
  pivot_longer(c(avg_gdpc, avg_cn, av_educ, urban), names_to = "variable", values_to = "value")
```

With this dataset, we can create a histogram with four facets:

Question x: Add the right geometry element to complete the histograms.



```
final_long %>%
  ggplot(aes(x = value, group = variable, fill = variable)) +
  facet_wrap(~variable, scales = "free") +
  # fill in the right answer
  geom_histogram()
```



Next, we would like to make a plot of some of the data. In particular, we can use the `ggplot` library to create a scatterplot of urbanization in 1500 on the x axis and \log (GDP per capita) on the y axis. That would allow us to get insight into the *reversal of fortune* theory.

Question: Use `ggplot` to create a plot as described above. Hint: use `geom_point()` as the geometric attribute. Also try to see if you can display the country name corresponding to each dot.

```
final %>%
  ggplot(aes(x = urban, y = log(avg_gdpc))) +
  geom_point() +
  geom_text(aes(label=country), nudge_x = 0.01, nudge_y = 0.2)
```



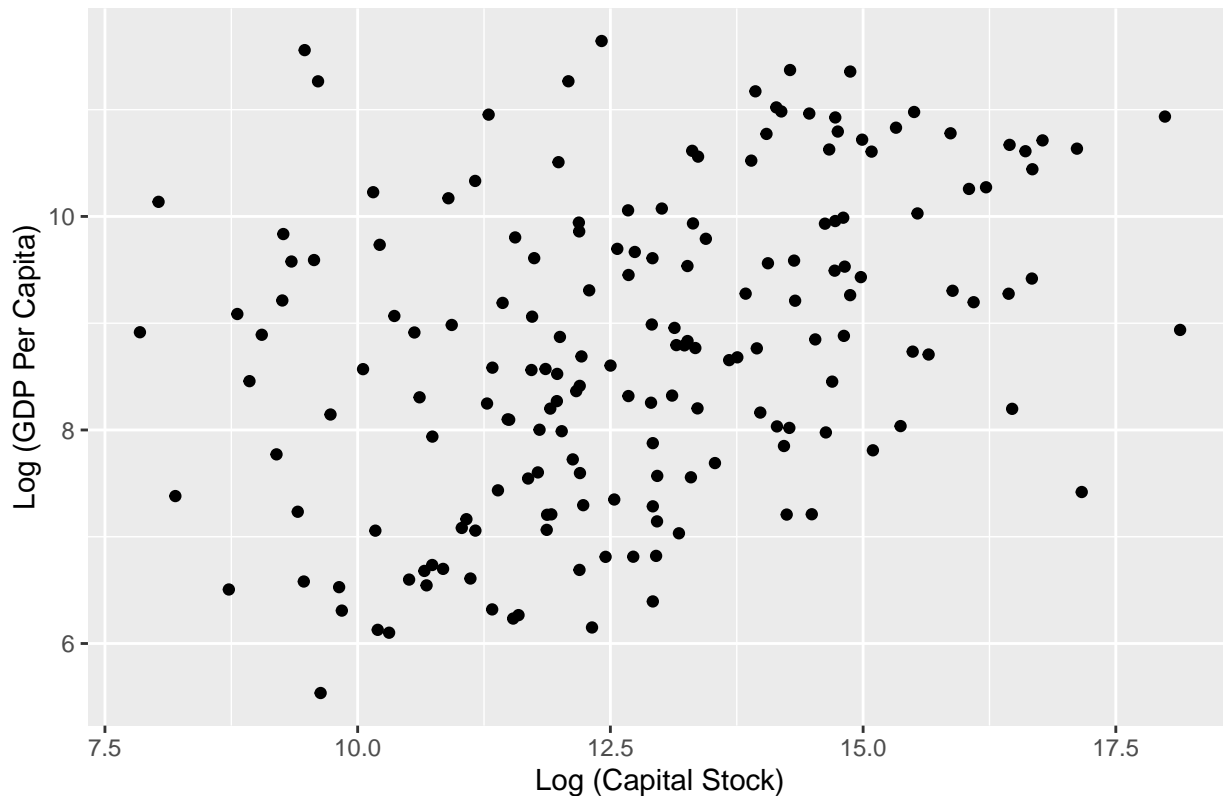
Question: Do you interpret this as evidence for the *reversal of fortune* theory? Why (not)?

Secondly, we can plot GDP per capita (y -axis) against the capital stock (x -axis). In this case, we want to log-transform both of the variables.

Question: Construct this plot. Also make sure to add a nice title and to change the axis titles appropriately.

```
final %>%
  ggplot(aes(x=log(avg_cn), y = log(avg_gdpc))) + geom_point() +
  xlab("Log (Capital Stock)") +
  ylab("Log (GDP Per Capita)") +
  ggtitle("Relationship between Capital Stock and Economic Well-Being")
```

Relationship between Capital Stock and Economic Well-Being



Question: Do you interpret this as evidence for the logic of most macroeconomic models, that a higher capital stock causes a higher level of income? Why (not)?

Finally, we would also like to make a map displaying economic growth rates. In order to do so, we need the `sf` package, short for *Spatial Features*. This is an efficient format in which data used to construct maps are stored. We also need a couple of auxiliary packages:

```
library(tidyverse); library(sf)
```

Possibly, we also have to install a couple of auxiliary packages:

```
pacman::p_load("rgdal", "rgeos", "lwgeom")
```

Let us first find a map of the world:

```
library(maps)
world <- st_as_sf(map("world", plot = FALSE, fill = TRUE))
```

We converted the world map from the `maps` package to an `sf` data.frame.

```
world
```

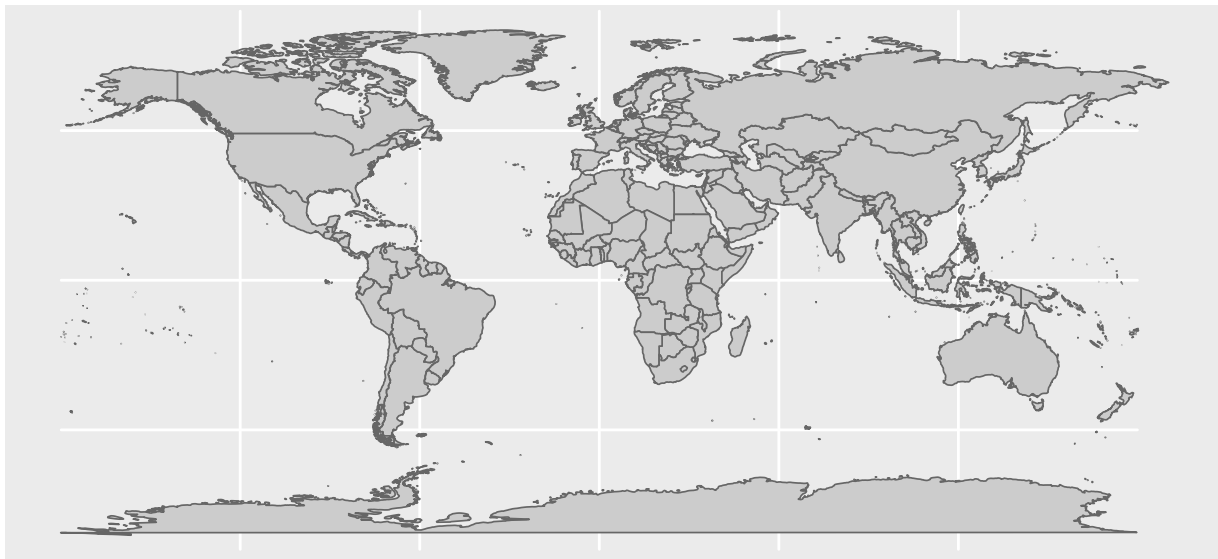
```
## Simple feature collection with 253 features and 1 field
## Geometry type: MULTIPOLYGON
## Dimension:      XY
## Bounding box:   xmin: -180 ymin: -85.19218 xmax: 190.2708 ymax: 83.59961
## Geodetic CRS:   WGS 84
## First 10 features:
##                ID                geom
```

```
## 1          Aruba MULTIPOLYGON (((-69.89912 1...
## 2    Afghanistan MULTIPOLYGON (((74.89131 37...
## 3          Angola MULTIPOLYGON (((23.9665 -10...
## 4    Anguilla MULTIPOLYGON (((-63.00122 1...
## 5          Albania MULTIPOLYGON (((20.06396 42...
## 6          Finland MULTIPOLYGON (((20.61133 60...
## 7          Andorra MULTIPOLYGON (((1.706055 42...
## 8 United Arab Emirates MULTIPOLYGON (((53.92783 24...
## 9          Argentina MULTIPOLYGON (((-64.54916 -...
## 10         Armenia MULTIPOLYGON (((45.55235 40...
```

world is now a data.frame, containing the names of countries and associated polygons. We can use this object to create a simple map:

```
world_map <- ggplot(world) +
  geom_sf(fill = "grey80",
          col = "grey40",
          lwd = 0.3)
```


world_map



It is also possible to change projections. [Here](#) is a short primer on different projections. For you, this is not particularly relevant, but it allows you to pick a projection which you like. Here's an example:

```
world_map +
  coord_sf(crs = "+proj=laea +y_0=0 +lon_0=0 +lat_0=0") +
  labs(subtitle = "Lambert Azimuthal Equal Area projection")
```


Lambert Azimuthal Equal Area projection



Should we want to zoom in on a particular part of the world, that is also possible. To do that, we can filter the dataframe based on many features, for example, on coordinates:

```
sf_use_s2(FALSE)

world %>%
  st_crop(xmin = -22,
          ymin = 35,
          xmax = 40,
          ymax = 70) %>%
  ggplot() + geom_sf()
```

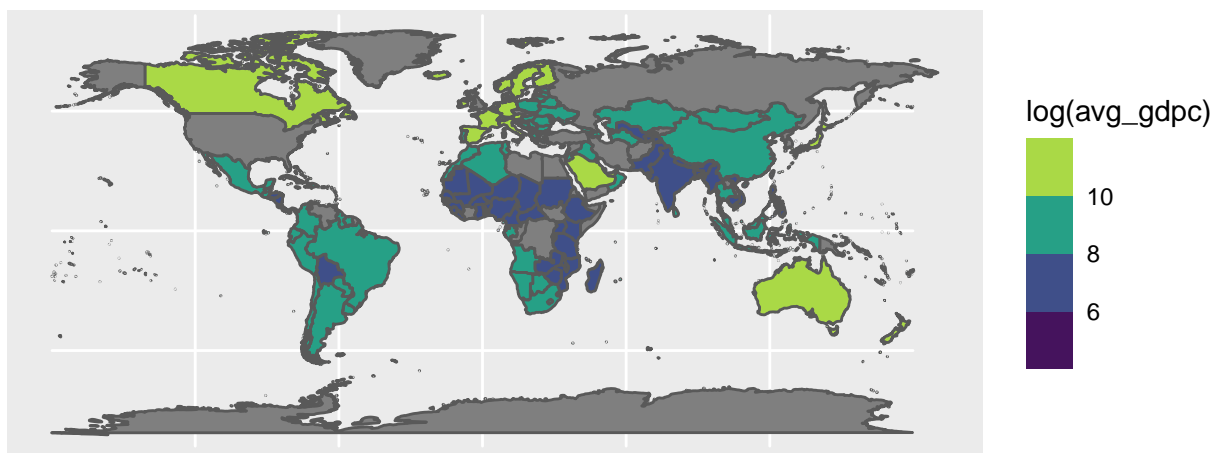


Question x: Take your `data.frame` `world`, and merge it with the final `data.frame` containing the GDP growth rates. Save this to a `data.frame` called `world_data`.

```
world_data <- world %>%
  left_join(final, by = c("ID" = "country"))
```

Now, plot a map displaying `avg_gdpc` in different countries. Hint: use the `geom_sf(fill = avg_gdpc)` as geometric attribute.

```
world_data %>%
  ggplot() + geom_sf(aes(fill = log(avg_gdpc))) + scale_fill_viridis_b()
```



The End